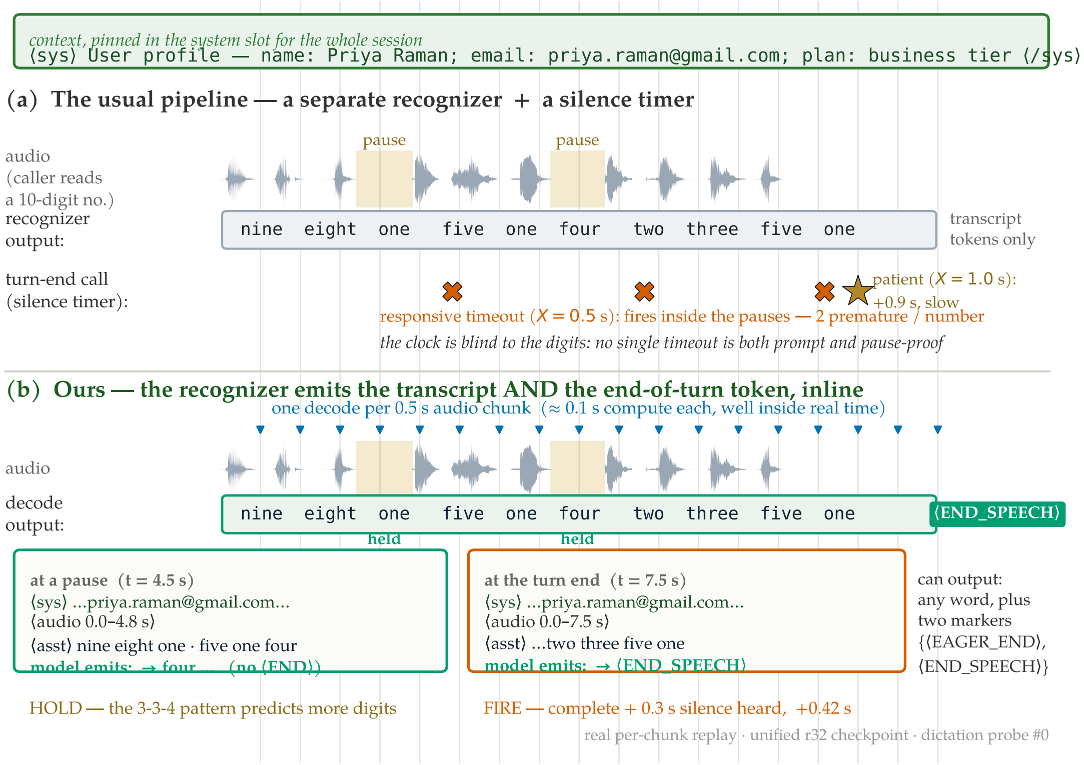
context, pinned in the system slot for the whole session
⟨sys⟩ User profile — name: Priya Raman; email: priya.raman@gmail.com; plan: business tier ⟨/sys⟩
(a) The usual pipeline — a separate recognizer + a silence timer
pause
pause
audio (caller reads a 10-digit no.)
recognizer output:
nine eight one five one four two three five one
transcript tokens only
turn-end call (silence timer):
patient (X = 1.0 s): +0.9 s, slow
responsive timeout (X = 0.5 s): fires inside the pauses — 2 premature / number
the clock is blind to the digits: no single timeout is both prompt and pause-proof
(b) Ours — the recognizer emits the transcript AND the end-of-turn token, inline
one decode per 0.5 s audio chunk (≈ 0.1 s compute each, well inside real time)
audio
decode output:
nine eight one five one four two three five one
⟨END_SPEECH⟩
held
held
at a pause (t = 4.5 s)
⟨sys⟩ ...priya.raman@gmail.com...
⟨audio 0.0–4.8 s⟩
⟨asst⟩ nine eight one · five one four
model emits: → four ... (no ⟨END⟩)
at the turn end (t = 7.5 s)
⟨sys⟩ ...priya.raman@gmail.com...
⟨audio 0.0–7.5 s⟩
⟨asst⟩ ...two three five one
model emits: → ⟨END_SPEECH⟩
can output: any word, plus two markers {⟨EAGER_END⟩, ⟨END_SPEECH⟩}
HOLD — the 3-3-4 pattern predicts more digits
FIRE — complete + 0.3 s silence heard, +0.42 s
real per-chunk replay · unified r32 checkpoint · dictation probe #0